

The Partnership Prioritization Framework

A practical method for deciding which partnerships deserve investment, which deserve monitoring, and which should quietly wind down.

BRIT PRINGLE • JULY 2026

01 The Question This Framework Answers

Most partnership teams don't fail because they pick the wrong partners. They fail because they try to work with too many of the right ones at once.

Every partner looks promising in the first meeting. Someone senior wants to explore product integration, or a BD rep signs a co-marketing agreement, and your team ends up managing a dozen active relationships without the systems to properly run two strategic ones. The impact of each real opportunity gets diluted by the operational load of servicing everyone.

That's the core mistake: treating partnership investment as a headcount problem instead of an allocation problem. More partners does not mean more results. It often means less, because attention is finite and every partner competes for the same executive time, engineering cycles, and go-to-market bandwidth.

The insight this framework encodes: differentiated investment produces differentiated returns. Not every partnership deserves the same resourcing, and pretending otherwise is how strong partnerships get starved while weak ones get overfunded. The job isn't to be equally nice to everyone. It's to decide, with evidence, where your next dollar and your best person's next hour should go.

This framework is a decision tool for that allocation problem. It gives you a repeatable way to score an opportunity, place it on a matrix that shows where it actually sits, and take the action that follows — invest, fix, delegate, or exit.

WHAT THIS FRAMEWORK WILL DO

- Give you a structured, repeatable way to score any partnership opportunity
- Force an honest read on execution feasibility, not just strategic appeal
- Tell you what action follows from a given score, so scoring isn't just an exercise
- Give you language for deprioritizing a partnership without burning the relationship

WHAT THIS FRAMEWORK WILL NOT DO

- Predict which partnerships will succeed — it reduces guesswork, it doesn't eliminate it
- Replace judgment about market timing, competitive dynamics, or the people involved
- Work as a one-time exercise — re-run it quarterly as circumstances change
- Substitute for the hard conversations that follow once you have a score

02 Partnership Tier Definitions

Three tiers, defined by depth of integration and strategic weight. Most portfolios are top-heavy with light-touch relationships and thin on transformational ones. This is usually correct, not a problem to fix.

One clarification before the definitions, because it prevents a common confusion: tiers describe how much you invest. They are not the same as the four partnership models in Section 3, which describe how a partnership actually creates value and moves money. The two are independent. For example, a value-added resale partner might sit in any tier depending on its strategic weight. Score the model; size the tier.

Tier 1 — Transactional Partners

WHAT IT LOOKS LIKE

Marketplace listing, reseller agreement, basic co-marketing (a webinar, a content swap). Contact is periodic, not constant. Mostly self-serve on both sides.

WHAT EACH SIDE GIVES AND RECEIVES

You give access to your marketplace or reseller terms and standard collateral. They give distribution reach and occasional deal referrals. Neither side customizes much for the other.

TYPICAL INVESTMENT

A fraction of one person's time. No dedicated engineering. No standing meeting cadence required.

HOW YOU KNOW IT'S WORKING

Referral volume or marketplace conversion stays positive without you working harder each quarter to get the same result. If output requires increasing effort to maintain, the relationship is quietly becoming a cost center, not a channel.

Tier 2 — Strategic Partners

WHAT IT LOOKS LIKE

A real co-selling motion — named reps on both sides, a joint go-to-market plan, recurring business reviews, shared pipeline visibility.

WHAT EACH SIDE GIVES AND RECEIVES

You commit dedicated BD or sales time and enablement. They commit field alignment and usually co-investment in campaigns or events. This is not free for either party.

TYPICAL INVESTMENT

A named owner (full-time at scale), cross-functional support from sales and marketing, executive visibility without necessarily weekly executive time.

HOW YOU KNOW IT'S WORKING

Attributable pipeline and revenue grow quarter over quarter, and the partner's field team can name your product unprompted. If your team is the only one showing up to business reviews, it's transactional wearing a strategic label. Downgrade it and stop pretending.

Tier 3 — Transformational Partners

WHAT IT LOOKS LIKE

Joint product integration, shared roadmap conversations, sometimes a co-branded or embedded offering. Engineering and product leadership are involved, not just BD.

WHAT EACH SIDE GIVES AND RECEIVES

You commit engineering resources, roadmap flexibility, and executive sponsorship. They commit the same, often plus some form of priority placement or exclusivity.

TYPICAL INVESTMENT

An executive sponsor on both sides, dedicated engineering allocation, a joint planning cadence that goes beyond quarterly check-ins. Expensive to build, expensive to maintain. This is why it should be rare.

HOW YOU KNOW IT'S WORKING

The integration is actually adopted by customers, not just announced. A press release is not evidence. Usage and renewal data tied to the integration is.

03 The Four Partnership Models

Tiers tell you how much to invest. Models tell you how a partnership actually creates value and moves money. Each of the four runs on different economics and fails in different ways. The organizing idea is simple: a partner can sell your product, influence its sale, extend it, or embed it.

The spine that runs through all four: most frameworks score whether a partner *can* execute. Almost none score whether the partner makes money when they lead with you — and that, not capability, is what predicts whether they actually *will*. Every model below lives or dies on that question.

SELL Value-Added Resale

The partner takes your product to market wrapped in delivery they own and the customer pays for: implementation, integration, IP, a vertical solution, or a managed service. The wrapper is the point. Strip it away and you have pure resell — the partner processes the PO, adds nothing the customer pays them for, and the margin races to zero. That's a billing relationship, not a partnership.

The test isn't "do they resell the license," it's "does the partner make money on their own value, or only on your discount?" Only the first is durable, because only the first gives the partner's field a reason to lead with you. (Referral and agency partners — pass a lead for a fee, no delivery — are the lightweight on-ramp to this model and to co-sell, not a model of their own.)

INFLUENCE Co-Sell

Joint selling with no resale: the partner sources or influences the deal, and the customer buys direct from you. This runs on field-level trust and a clean attribution model, not margin.

The failure mode is specific: an executive-signed co-sell agreement produces nothing until individual reps on both sides believe the other will show up in their specific deal and that trust is rebuilt account by account, not announced once. Partner economics still governs. The overlay rep who gets comped on your deal shows up; the one who doesn't, won't.

EXTEND Product Integration

A technical integration that extends value for mutual customers. The currency is adoption and influenced renewal, where usage and renewal data is tied to the integration, never just a press release.

The economics question here is subtler but no less real: an integration only gets sold if someone owns selling it after it ships. The most common death is the handoff gap — engineering builds it, no one owns the go-to-market, and it goes live but dark. Score whether there is a revenue owner for the integration, not just a build owner.

EMBED OEM / White-Label

Your product is embedded inside the partner's product and sold under their brand or as a component; the end customer may never know you exist. It's economically distinct from resale in every way that counts, such as transfer pricing instead of discount off list, rev-share and volume commitments, and a deep dependency, because the partner's roadmap and their churn become your revenue.

The partner's economics are strongest here — your capability becomes part of what they sell — but so is your exposure, which is why embed deals warrant the heaviest diligence on the partner's own trajectory.

A note on marketplace. Cloud marketplace transactions — AWS, Azure, and GCP private offers and co-sell programs — are a channel that runs *through* resale and co-sell, not a fifth model. Treat marketplace as connective tissue for how these models transact, not a category of its own.

Which model you're in changes what a strong score looks like. "Revenue Potential" below means direct bookings for resale, influenced renewal for integration, and rev-share against a volume commitment for embed. Read the rubric with your model in mind.

04 The Scoring Rubric

Score every opportunity against seven criteria, each weighted by importance. Multiply each 1–5 score by its weight, sum the results, and you get a total out of 100. Formula: weighted score = (raw score ÷ 5) × weight, summed across all seven.

Revenue Potential 25%

Realistic 12-month revenue impact if the partnership succeeds.

What "revenue" means depends on your model (Section 3): direct bookings for resale, influenced renewal for integration, rev-share for embed.

1 LOOKS LIKE

No credible path to revenue in 12 months, or the estimate is a guess with no pipeline, customer signal, or comparable deal behind it.

5 LOOKS LIKE

A defensible number backed by a named pipeline, a signed pilot, or a closely comparable partnership's actual results.

ASK — *Where does this number come from? If you strip out the most optimistic 20%, does it still justify the investment?*

Partner Economics 15%

Does the partner make money — margin, services, retention, differentiation — when they lead with you, in a way their field can see?

This is the predictor. Partner Readiness asks whether they can execute; Partner Economics asks whether they have a reason to.

1 LOOKS LIKE

The only thing in it for the partner is your discount or a vague "strategic" halo. Their field has no dollar reason to prioritize you over the other vendors already in their stack.

5 LOOKS LIKE

The partner's own P&L improves when they push you — services attach, higher retention, a differentiated offer — and their reps can see the upside in their own numbers, not just their BD team's slides.

ASK — *When a partner rep chooses you over an alternative in a live deal, what do they personally get? Can they see it in their comp or quota, or only in a slideware "alignment"?*

Strategic Fit 15%

Does this align with your own product and market direction, not just offer a convenient swap.

1 LOOKS LIKE

The partner's audience or roadmap has little overlap with where you're actually headed in the next 12–18 months.

5 LOOKS LIKE

The partnership sits on your existing roadmap and go-to-market plan — you'd pursue this direction even without this specific partner.

ASK — *Would we build toward this outcome anyway? Does this pull us toward our strategy or away from it?*

Execution Feasibility 15%

Can you actually deliver on your side without breaking other commitments.

1 LOOKS LIKE

Delivering requires headcount or engineering time you don't have, and getting it means cutting something that matters more.

5 LOOKS LIKE

You can deliver with current resourcing, or the incremental ask is small and clearly fundable.

ASK — *What has to be true, or who has to say yes, for us to deliver our side? What are we not doing if we say yes to this?*

Ecosystem Leverage 10%

Does this partnership open doors to other partnerships or customers.

1 LOOKS LIKE

Self-contained. Winning it doesn't help you win anything else.

5 LOOKS LIKE

A credible unlock — access to a category, a position, or downstream relationships you couldn't easily get otherwise.

ASK — *If this works, what becomes easier that wasn't easier before? Is there a compounding effect, or is this a one-off?*

Partner Readiness 10%

Does the partner have the field, product, and executive muscle to execute their side.

1 LOOKS LIKE

Interest without dedicated resourcing, no clear internal champion, or a track record of commitments that didn't materialize.

5 LOOKS LIKE

A named owner, budget, and a prior execution history with partners like you.

ASK — *Who specifically owns this on their side, and what's their track record? Have they delivered on something similar with someone else?*

Opportunity Cost 10%

What you would not do if you invested here.

1 LOOKS LIKE

Pursuing this means deprioritizing a higher-confidence opportunity, and nobody has said that tradeoff out loud.

5 LOOKS LIKE

This doesn't meaningfully compete with your other priorities for the same scarce resource — same sponsor, same team, same quarter.

ASK — *What's the next-best use of the time and people this requires? If you could pursue only one partnership this quarter, would it be this one?*

WORKED EXAMPLE

A mid-market integration opportunity scores Revenue Potential 3, Partner Economics 3, Strategic Fit 4, Execution Feasibility 2, Ecosystem Leverage 3, Partner Readiness 4, Opportunity Cost 3.

Weighted: $15 + 9 + 12 + 6 + 6 + 8 + 6 = 62/100$. That lands in the "hold and reassess" band from Section 6 — not a strategic-tier partnership yet, but worth watching.

05 The Prioritization Matrix

Once you've scored an opportunity, plot it. A number alone doesn't tell you what to do, placement does.

Two composite axes. Strategic weight (Y-axis) = (Strategic Fit + Ecosystem Leverage) ÷ 2.
 Execution reality (X-axis) = (Execution Feasibility + Partner Readiness + Partner Economics) ÷ 3.
 Both run 1–5; split each axis at the midpoint to get four quadrants. Partner Economics belongs on the execution axis on purpose: a partner who *can* deliver but has no reason to isn't execution reality, it's wishful thinking.

↑ Strategic weight	Fix or Defer Real opportunity, can't deliver yet. Don't invest more until execution becomes real.	Invest Deeply Your strategic-tier partnerships. Full resourcing, executive sponsor, standing cadence.
	Exit or Ignore Stop pretending. The quadrant most teams avoid looking at honestly.	Automate or Delegate You can deliver, but it doesn't move anything that matters. Systematize it.

Execution reality →

A note on precision: the matrix is directional, not exact. Don't let two partnerships fight over a tenth of a point. If they're close, they're functionally in the same place. Use it to spot clusters and outliers, not to rank-order your portfolio to two decimal places.

o6 The Decision Framework

Your total score maps to a default action. Defaults can be overridden with judgment, but you should have to argue your way out of them, not into them.

80+	Strategic tier. Full investment, quarterly business review, named executive sponsor.
60–79	Hold at current tier. Revisit for strategic promotion after six months of demonstrated execution, not promised execution.
40–59	Transactional tier. Automate what you can. Minimum viable relationship — keep it alive, don't feed it.
Below 40	Exit path within 90 days.

COMMUNICATING DEPRIORITIZATION WITHOUT BURNING THE RELATIONSHIP

This is the part most frameworks skip, and it's the part that actually determines whether the framework is usable.

- **Be honest about the reason, not just the decision.** "We're not seeing the execution capacity to make this a priority right now" is truer and kinder than a vague brush-off, and it gives the partner something specific to fix.
- **Don't ghost.** A partner who gets silence instead of a decision will assume the worst about you, and they'll be right to. A short, direct conversation costs twenty minutes and protects your reputation in a market smaller than it looks.
- **Offer the honest downgrade, not a fake maintain.** If a partnership is dropping a tier, say so plainly and describe what the relationship looks like now. Don't pretend nothing has changed while quietly pulling resources.
- **Leave a real door open.** Circumstances change. A partner who scores low this year might score much higher next year if their execution catches up. Say so, specifically.

07 What This Framework Doesn't Do

- **It does not model long-term ecosystem dynamics.** Some partnerships are worth keeping for optionality five years out, not returns this year. This framework doesn't see that far.
- **It does not account for personal relationships that can override the math.** Sometimes a partnership survives because of one relationship at the top, and that's a real dynamic, not a flaw in the world. Just don't confuse it with strategy.
- **It does not replace judgment about market timing or competitive shifts.** A partnership that scores low today might be exactly the right bet if the market is about to move. The score is an input to judgment, not a substitute for it.
- **It does not evaluate consumption-based partnership economics.** Usage-based models run on different math than the revenue-and-effort logic this framework assumes. That's a different tool.

A scoring model isn't the hard part of this work. The hard part is deciding what question you're actually asking — not "which partners are loudest right now" but "which of these relationships actually deserves my team's time." Everything in this framework exists to make that second question answerable instead of just debatable.

ADDENDUM Where AI Fits Into This

Everything in Section 4 is mechanical enough to partially automate. That's not an accident. A weighted rubric with defined criteria, a threshold, and an action that follows from the score is exactly the kind of structured judgment a language model can help with, as long as you keep it in the right place.

The pattern is simple: give a model your rubric in plain language, hand it one input at a time — a signal, a partner data point, whatever you're triaging — and ask for a structured score plus a short recommended action, the same shape as Section 4, just applied at higher volume than a person can sustain. Set a numeric threshold. Anything below it gets logged, not acted on. Anything above it lands in front of a human.

I built a small version of this for myself after leaving AWS. It watches public signals about companies I'd want to partner with and scores each one the same way this framework scores a partnership opportunity: defined criteria, multiplied into a composite, checked against a threshold. It's a monitoring layer, not a decision-maker.

That distinction matters. Treat the model like a very literal junior analyst: it will apply your criteria exactly as written, which means the criteria have to be genuinely well-specified, and you have to check its work until you trust it. The first pass is never the last pass. You'll spend real time correcting its judgment before it earns the right to run with less supervision. That's not a flaw in the approach. It's the actual cost of automating a triage step, and it's worth paying because the alternative is a human doing volume work that a model does more consistently.

What it doesn't change: the framework still makes the call on what to do with a strategic partnership. AI can clear the noise so your judgment gets spent on the fraction of signals that deserve it. It doesn't replace the judgment itself, because that was never up for automation.